

# RePaint: Inpainting using Denoising Diffusion Probabilistic Models

Kate Liang, Huining Liu, Hongrui Tang, Jennie Wu, Mengjia Zhao

Date: May 1st, 2025

## BACKGROUND & MOTIVATION

### Image Inpainting

- Filling missing regions within an image




Figure 1: Diffusion-based inpainting

### Context & Motivation

- Existing AR and GAN methods struggle with large or sparse masks.
- Diffusion models lacked harmonization of the known and generated part.

### Goal

- Implement the reverse diffusion + resampling approach introduced by the paper
- Reproduce LPIPS results on the Celeb-A-HQ dataset

## DATA & MODEL

- Dataset:** CelebA-HQ-256
- Masks:** Wide, Narrow, Half, Expand, Alternate Lines, Super-Resolution
- Model:** Pretrained google/ddpm-celebahq-256 DDPM from Hugging Face

## METHODS




Figure 2: Overview of the approach


- Condition Inference:** Generate image by guiding a pre-trained diffusion process conditioned on masked input
- No Model Retraining:** Leverage pre-trained diffusion model
- Content Harmonization:** Harmonize known and generated regions through iterative resampling
- Bidirectional diffusion steps:** Alternate forward and backward in time
- Larger jumps in diffusion time:** Improve perceptual quality and diversity

## EXPERIMENT RESULTS

Table 1: CelebA-HQ Quantitative Results. Comparison against state-of-the-art methods. We compute the LPIPS (lower is better) for 6 different mask settings.

CelebA-HQ	Wide LPIPS↓	Narrow LPIPS↓	Super-Res 2x LPIPS↓	Altern. Lines LPIPS↓	Half LPIPS↓	Expand LPIPS↓
AOT [6]	0.104	0.047	0.714	0.667	0.287	0.604
DSI [2]	0.067	0.038	0.128	0.049	0.211	0.487
ICT [4]	0.063	0.036	0.483	0.353	0.166	0.432
DeepFillv2 [5]	0.066	0.049	0.119	0.049	0.209	0.467
LaMa [3]	0.045	0.028	0.177	0.083	0.138	0.342
RePaint [1]	0.059	0.028	0.029	0.009	0.165	0.435
Ours	0.075	0.043	0.057	0.048	0.193	0.534

## ABLATION EXPERIMENT



### Design Choice:

- Utilized pre-trained model from HuggingFace
- Set total number of reverse steps to 1000
- Computed average LPIPS score for each mask based on 20 val images

### Algorithm 1 Inpainting using our RePaint approach.

```

1:  $x_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:   for  $u = 1, \dots, U$  do
4:      $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\epsilon = \mathbf{0}$ 
5:      $x_{t-1}^{\text{known}} = \sqrt{\alpha_t} x_0 + (1 - \bar{\alpha}_t) \epsilon$ 
6:      $z \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $z = \mathbf{0}$ 
7:      $x_{t-1}^{\text{unknown}} = \frac{1}{\sqrt{\alpha_t}} \left( x_t - \frac{\beta_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right) + \sigma_t z$ 
8:      $x_{t-1} = m \odot x_{t-1}^{\text{known}} + (1 - m) \odot x_{t-1}^{\text{unknown}}$ 
9:     if  $u < U$  and  $t > 1$  then
10:       $x_t \sim \mathcal{N}(\sqrt{1 - \beta_{t-1}} x_{t-1}, \beta_{t-1} \mathbf{I})$ 
11:    end if
12:  end for
13: end for
14: return  $x_0$ 

```

Figure 3: RePaint Algorithm

## VISUAL RESULTS




Figure 4: Comparison against the state-of-the-art methods

## CONCLUSION & FUTURE WORK

**RePaint** is a novel conditioning method, with a unified inference schedule that works for a variety of mask type and diffusion models. It produces semantically coherent completions by harmonizing known and generated regions.

We were able to re-implement and mostly reproduce the CelebA-HQ results of the original paper.

### Future Work:

- Conditioning on language instruction and known area
- Image Inpainting → Video Inpainting